



CS598: VISUAL INFORMATION RETRIEVAL

Lecture III: Image Representation:

- Invariant Local Image Descriptors

RECAP OF LECTURE II

- Color, texture, descriptors
 - Color histogram
 - Color correlogram
 - LBP descriptors
 - Histogram of oriented gradient
 - Spatial pyramid matching
- Distance & Similarity measure
 - L_p distances
 - Chi-Square distances
 - KL distances
 - EMD distances
 - Histogram intersection





LECTURE III: PART I

Local Feature Detector



OUTLINE

- Blob detection
 - Brief of Gaussian filter
 - Scale selection
 - Lapacian of Gaussian (LoG) detector
 - Difference of Gaussian (DoG) detector
 - Affine co-variant region



GAUSSIAN KERNEL

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$




0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

5 x 5, $\sigma = 1$


- Constant factor at front makes volume sum to 1 (can be ignored when computing the filter values, as we should renormalize weights to sum to 1 in any case)

GAUSSIAN KERNEL

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



$\sigma = 2$ with 30 x 30
kernel





$\sigma = 5$ with 30 x 30
kernel

- Standard deviation σ : determines extent of smoothing


CHOOSING KERNEL WIDTH

- The Gaussian function has infinite support, but discrete filters use finite kernels




CHOOSING KERNEL WIDTH

- Rule of thumb: set filter half-width to about 3σ



GAUSSIAN VS. BOX FILTERING



GAUSSIAN FILTERS

- Remove “high-frequency” components from the image (low-pass filter)
- Convolution with self is another Gaussian
 - So can smooth with small- σ kernel, repeat, and get same result as larger- σ kernel would have
 - Convoluting two times with Gaussian kernel with std. dev. σ is same as convoluting once with kernel with std. dev. $\sigma\sqrt{2}$
- *Separable* kernel
 - Factors into product of two 1D Gaussians



SEPARABILITY OF THE GAUSSIAN FILTER

$$\begin{aligned} G_{\sigma}(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian



SEPARABILITY EXAMPLE

2D convolution
(center location
only)

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix}$$

The filter factors
into a product of
1D
filters:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

Perform
convolution
along rows:

$$\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix} = \begin{bmatrix} & 11 & \\ & 18 & \\ & 18 & \end{bmatrix}$$

Followed by convolution
along the remaining
column:



WHY IS SEPARABILITY USEFUL?

- What is the complexity of filtering an $n \times n$ image with an $m \times m$ kernel?
 - $O(n^2 m^2)$
- What if the kernel is separable?
 - $O(n^2 m)$





OUTLINE

- Blob detection
 - Brief of Gaussian filter
 - Scale selection
 - Lapacian of Gaussian (LoG) detector
 - Difference of Gaussian (DoG) detector
 - Affine co-variant region



BLOB DETECTION IN 2D

- Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D





$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$



BLOB DETECTION IN 2D

- Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D




Scale-normalized:

$$\nabla_{\text{norm}}^2 g = \sigma^2 \left(\frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$$




SCALE SELECTION


- At what scale does the Laplacian achieve a maximum response to a binary circle of radius r ?



image



Laplacian




SCALE SELECTION

- At what scale does the Laplacian achieve a maximum response to a binary circle of radius r ?
- To get maximum response, the zeros of the Laplacian have to be aligned with the circle

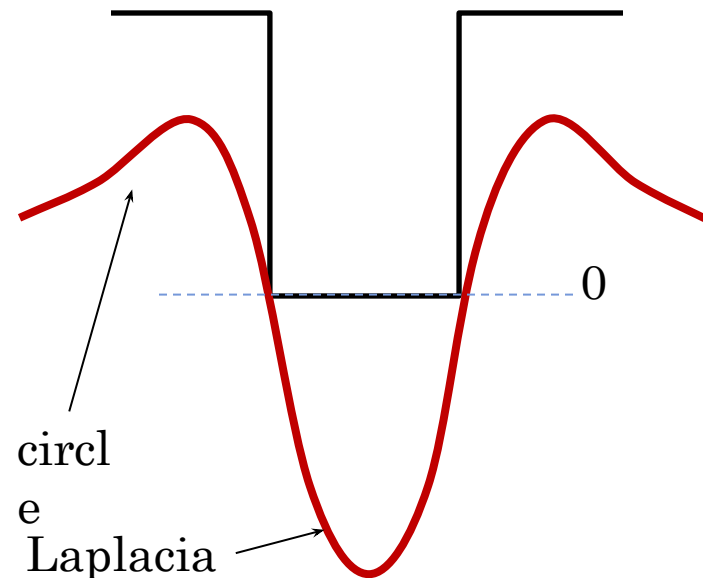
- The Laplacian is given by:

$$(x^2 + y^2 - 2\sigma^2) e^{-(x^2 + y^2)/2\sigma^2} / 2\pi\sigma^6$$

- Therefore, the maximum response occurs at $\sigma = r / \sqrt{2}$.



image




circle

Laplacian



CHARACTERISTIC SCALE

- We define the characteristic scale of a blob as the scale that produces peak of Laplacian response in the blob center



T. Lindeberg (1998). "Feature detection with automatic scale selection."
International Journal of Computer Vision **30** (2): pp 77--116.

SCALE-SPACE BLOB DETECTOR

1. Convolve image with scale-normalized Laplacian at several scales



SCALE-SPACE BLOB DETECTOR: EXAMPLE



SCALE-SPACE BLOB DETECTOR: EXAMPLE




sigma = 11.9912




SCALE-SPACE BLOB DETECTOR

1. Convolve image with scale-normalized Laplacian at several scales
2. Find maxima of squared Laplacian response in scale-space



SCALE-SPACE BLOB DETECTOR: EXAMPLE



OUTLINE

- Blob detection
 - Brief of Gaussian filter
 - Scale selection
 - Lapacian of Gaussian (LoG) detector
 - Difference of Gaussian (DoG) detector
 - Affine co-variant region



EFFICIENT IMPLEMENTATION


- Approximating the Laplacian with a difference of Gaussians:

$$L = \sigma^2 \left(G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma) \right)$$


(Laplacian)

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

(Difference of Gaussians)



EFFICIENT IMPLEMENTATION



David G. Lowe. **"Distinctive image features from scale-invariant keypoints."** *IJCV* 60 (2), pp. 91-110, 2004.

INVARIANCE AND COVARIANCE PROPERTIES

- Laplacian (blob) response is *invariant* w.r.t. rotation and scaling
- Blob location and scale is *covariant* w.r.t. rotation and scaling
- What about intensity change?




OUTLINE

- Blob detection
 - Brief of Gaussian filter
 - Scale selection
 - Lapacian of Gaussian (LoG) detector
 - Difference of Gaussian (DoG) detector
 - Affine co-variant region



ACHIEVING AFFINE COVARIANCE

- Affine transformation approximates viewpoint changes for roughly planar objects and roughly orthographic cameras




ACHIEVING AFFINE COVARIANCE

Consider the second moment matrix of the window containing the blob:

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$

Recall:


$$\begin{bmatrix} u & v \end{bmatrix} M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$



This ellipse visualizes the “characteristic shape” of the window



AFFINE ADAPTATION EXAMPLE



Scale-invariant regions
(blobs)



AFFINE ADAPTATION EXAMPLE





Affine-adapted
blobs




FROM COVARIANT DETECTION TO INVARIANT DESCRIPTION

- Geometrically transformed versions of the same neighborhood will give rise to regions that are related by the same transformation
- What to do if we want to compare the appearance of these image regions?
 - *Normalization*: transform these regions into same-size circles



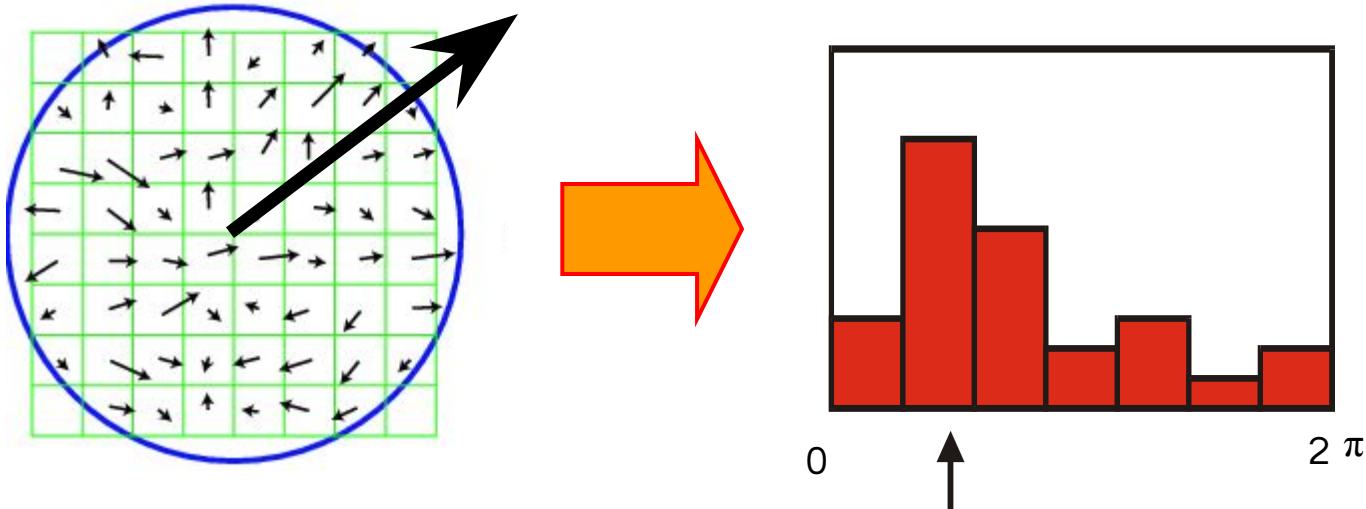
AFFINE NORMALIZATION

- Problem: There is no unique transformation from an ellipse to a unit circle
 - We can rotate or flip a unit circle, and it still stays a unit circle



ELIMINATING ROTATION AMBIGUITY

- To assign a unique orientation to circular image windows:
 - Create histogram of local gradient directions in the patch
 - Assign canonical orientation at peak of smoothed histogram



FROM COVARIANT REGIONS TO INVARIANT FEATURES

Extract affine
regions




Normalize
regions




Eliminate
rotational
ambiguity



Compute
appearance



SIFT (Lowe



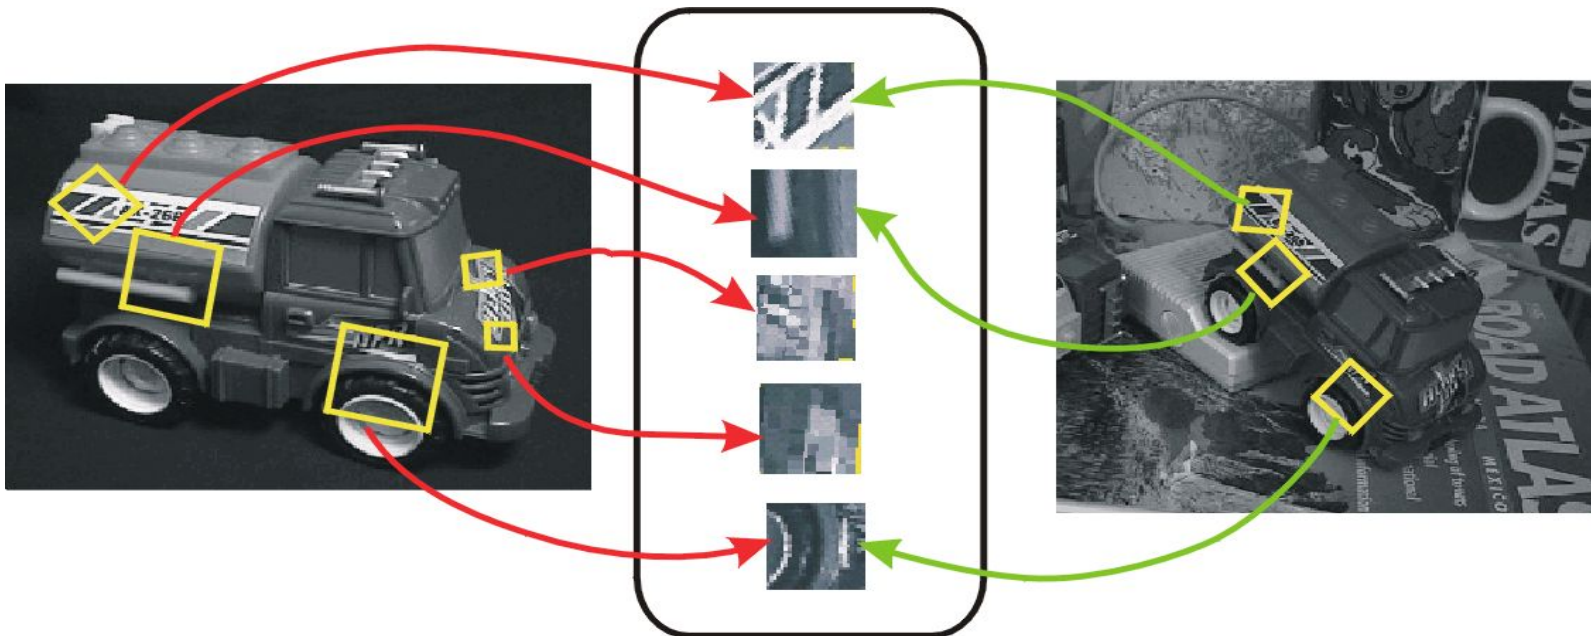
INVARIANCE VS. COVARIANCE

□ Invariance:

- $\text{features}(\text{transform}(\text{image})) = \text{features}(\text{image})$

□ Covariance:

- $\text{features}(\text{transform}(\text{image})) = \text{transform}(\text{features}(\text{image}))$



Covariant detection => invariant description



David G. Lowe, “*Distinctive Image Features from Scale-Invariant Keypoints*”, International Journal of Computer Vision, Vol. 60, No. 2, pp. 91-110

- 6291 as of 02/28/2010; 22481 as of 02/03/2014
- Our goal is to design the best local image descriptors in the world.

LECTURE III: PART I

Learning Local Feature Descriptor



Panoramic stitching [Brown et al. CVPR05]



Real-world Face recognition
[Wright & Hua CVPR09]
[Hua & Akbarzadeh ICCV09]



Image databases
[Mikolajczyk & Schmid ICCV01]




Object or location recognition
[Nister & Stewenius CVPR06]
[Schindler et al. CVPR07]



Courtesy of Seitz & Szeliski

Robot navigation
[Deans et al.]



3D reconstruction
[Snavely et al. SIGGRAPH06]

TYPICAL MATCHING PROCESS

Image 1



Image 2



Interest point/region detection
(sparse)

Image 1





Image 2



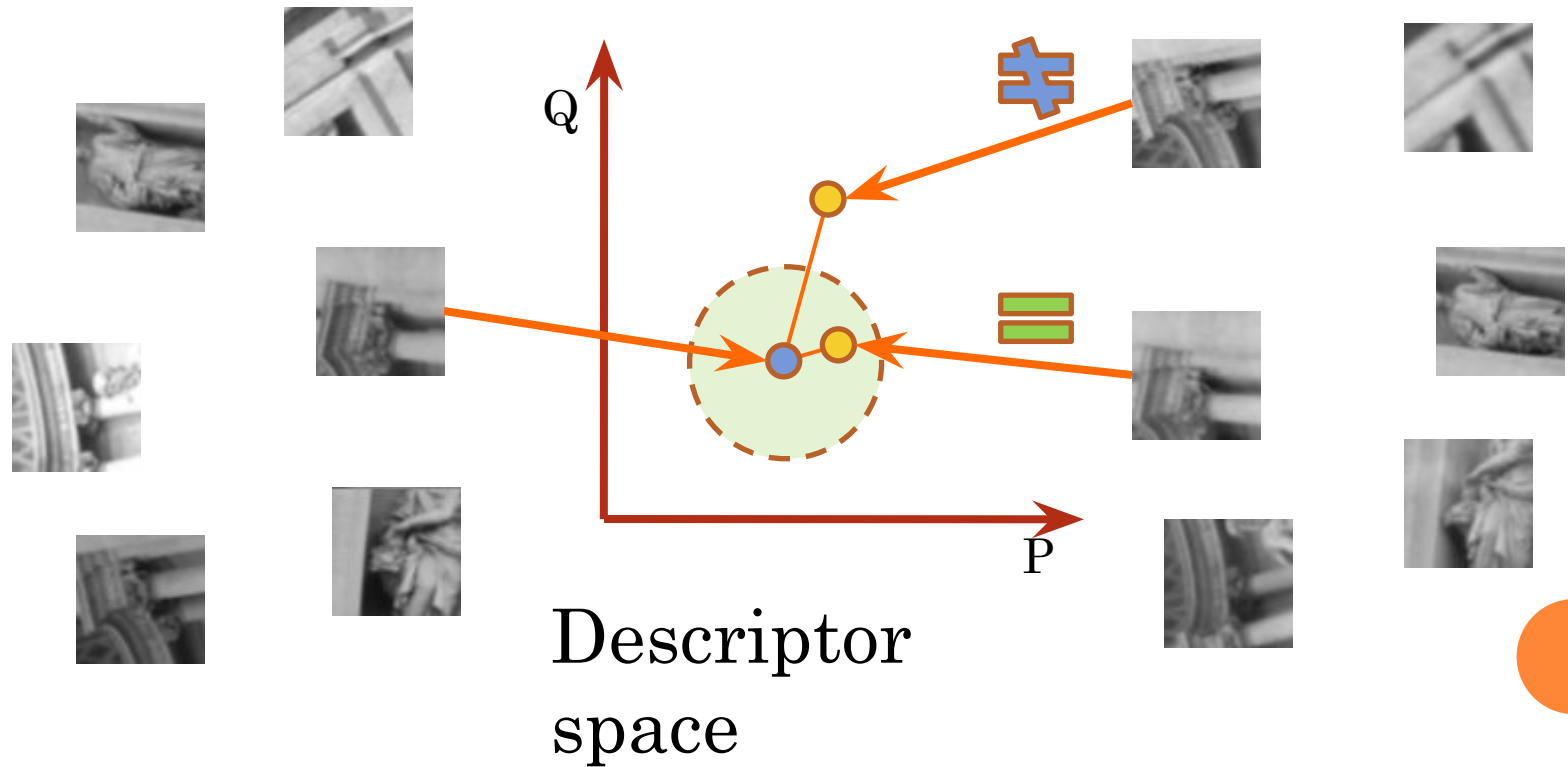
Dense sampling of image patches



TYPICAL MATCHING PROCESS

Image 1

Image 2




PROBLEM TO SOLVE

Learning a function of a local image patch
 $descriptor = f(\text{image patch})$
s.t. a nearest neighbor classifier is optimal


- To obtain the most *discriminative*, *compact*, and *computationally efficient* local image descriptors.
 - How can we get ground truth data?
 - What is the form of the descriptor function $f(\cdot)$?
 - What is the measure for optimality?



HOW CAN WE GET GROUND TRUTH DATA?



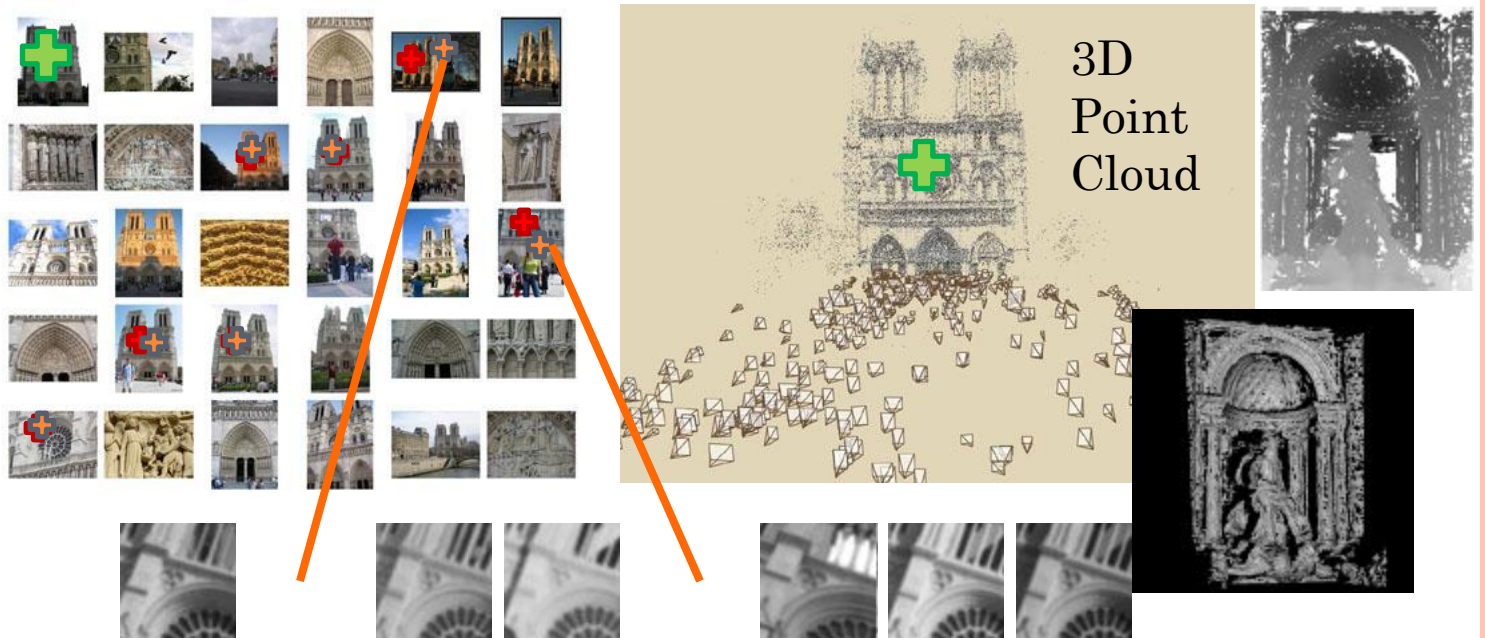
TRAINING DATA



Multiview stereo = Training data

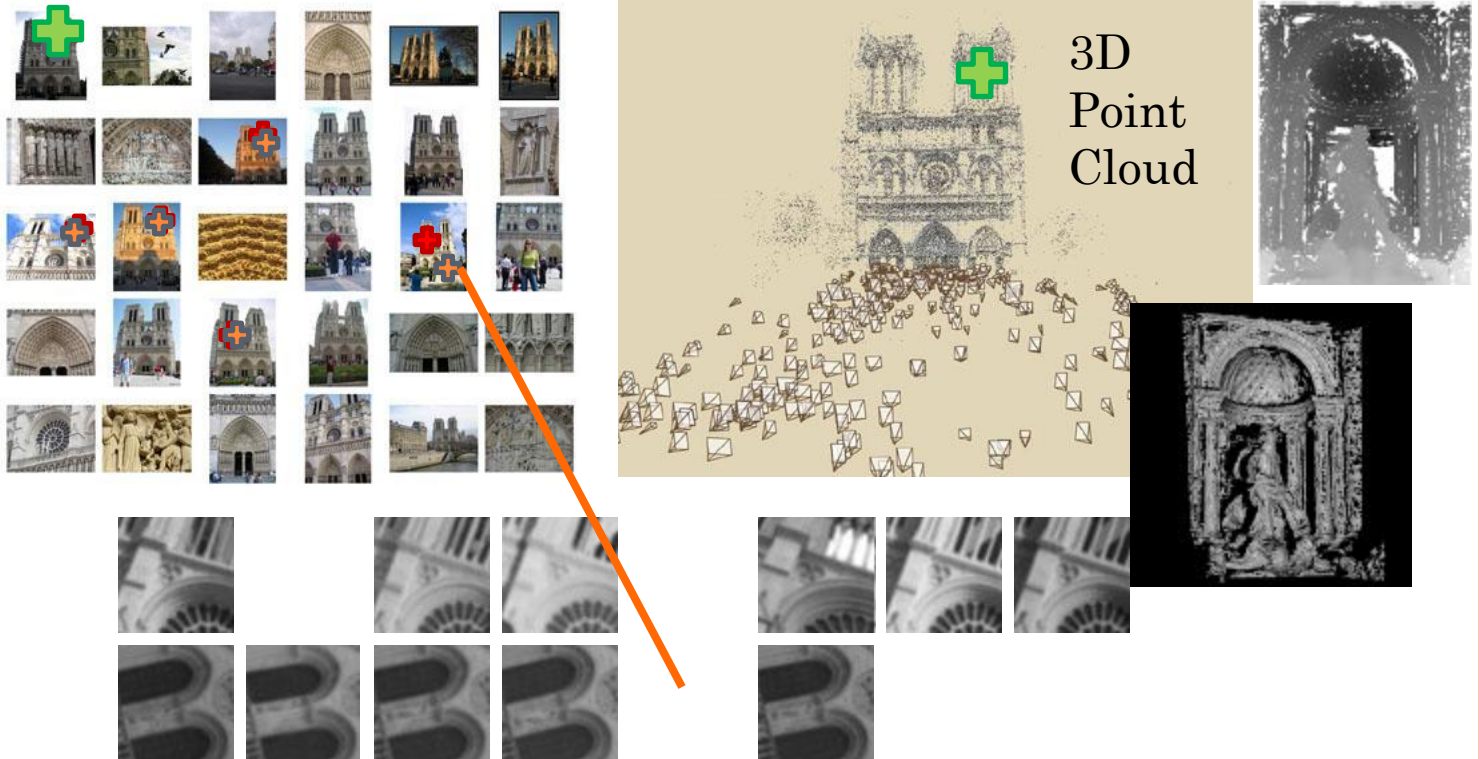
[Goesele et al. - ICCV'07] [Snavely et al. - SIGGRAPH'06]

TRAINING DATA



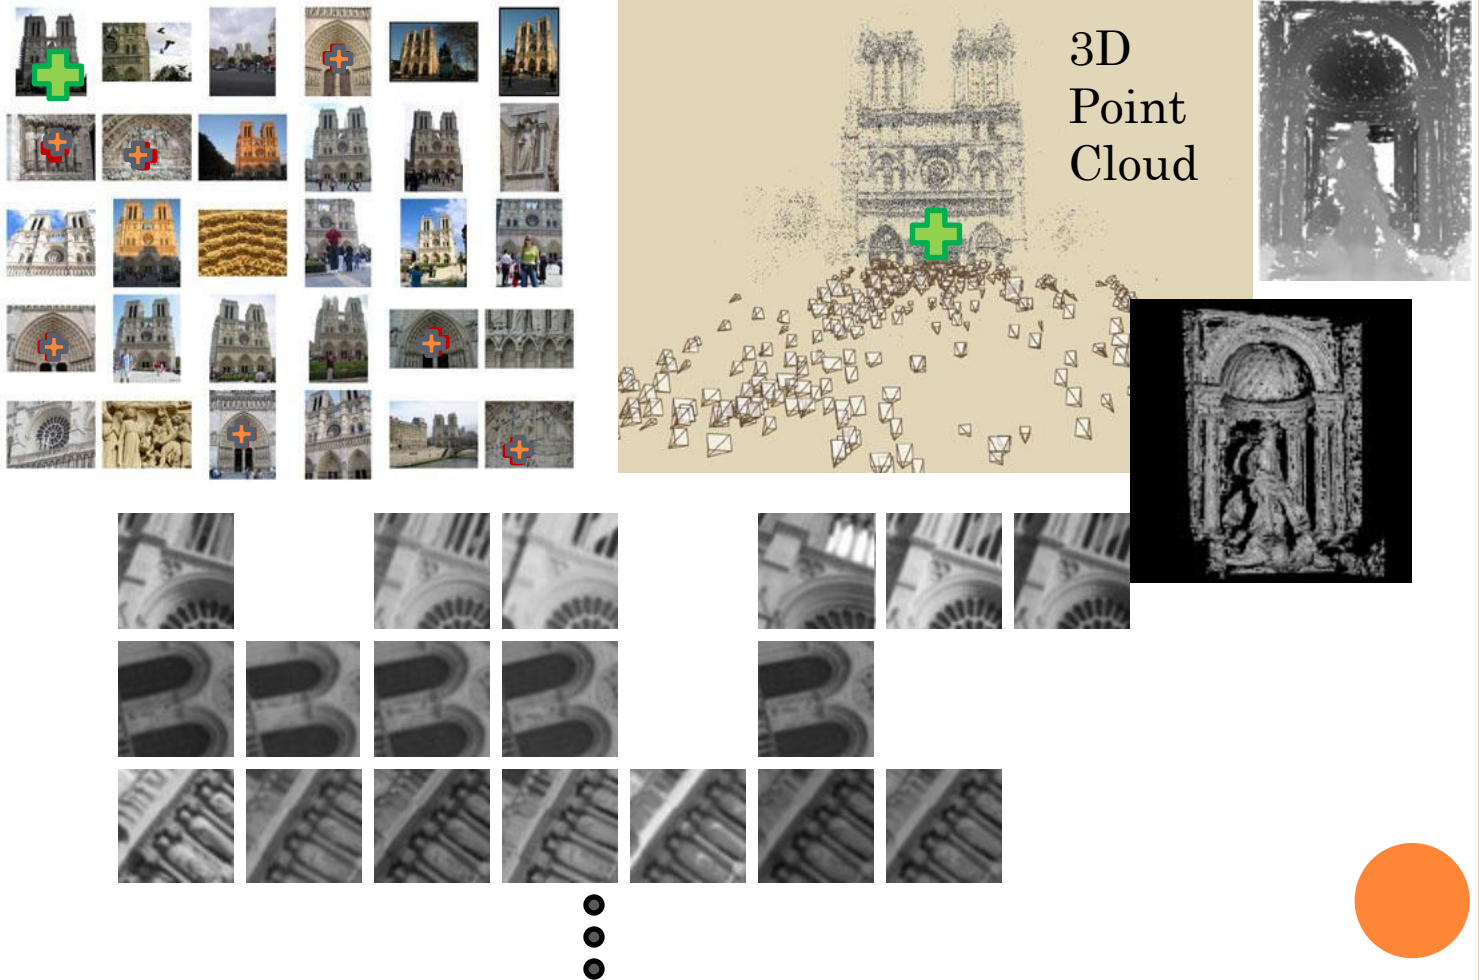
[Goesele et al. - ICCV'07] [Snavely et al. - SIGGRAPH'06]

TRAINING DATA




[Goesele et al. - ICCV'07] [Snavely et al. - SIGGRAPH'06]

TRAINING DATA




[Goesele et al. - ICCV'07] [Snavely et al. - SIGGRAPH'06]





Libert




Yosemit


e

- Statue of liberty (New York) – Liberty
- Notre Dame (Paris) – Notre Dame
- Half Dome (Yosemite) - Yosemite
- <http://www.cs.ubc.ca/~mbrown/patchdata/patchdata.html>

WHAT IS THE FORM OF THE DESCRIPTOR FUNCTION?





DESCRIPTOR ALGORITHMS



[SIFT – Lowe
ICCV'99]




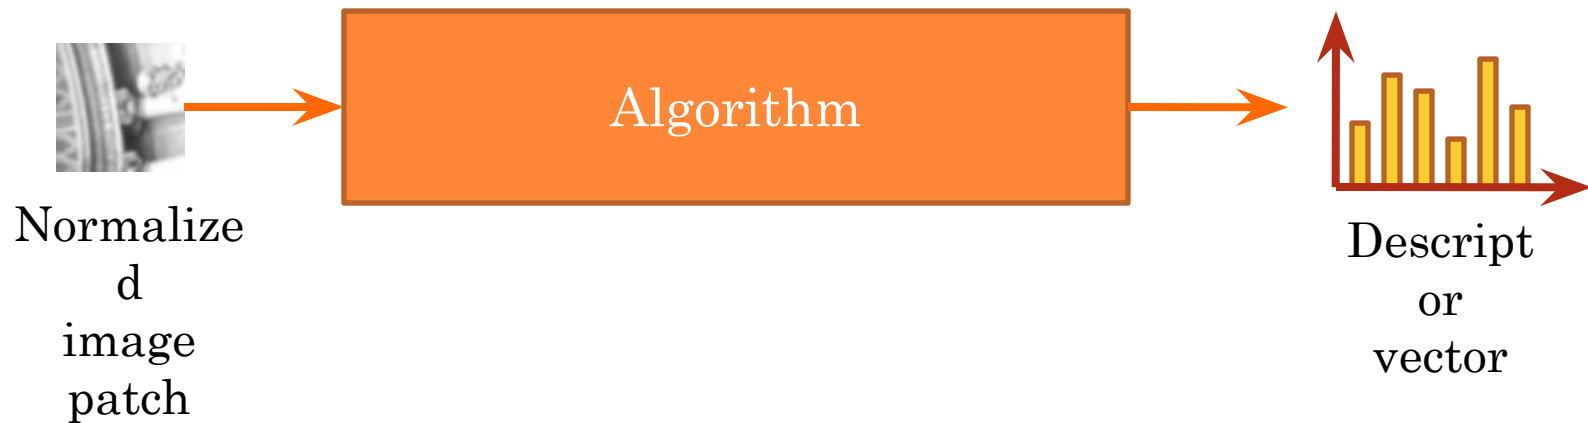
DESCRIPTOR ALGORITHMS



[GLOH – Mikolajczyk & Schmid
PAMI'05]



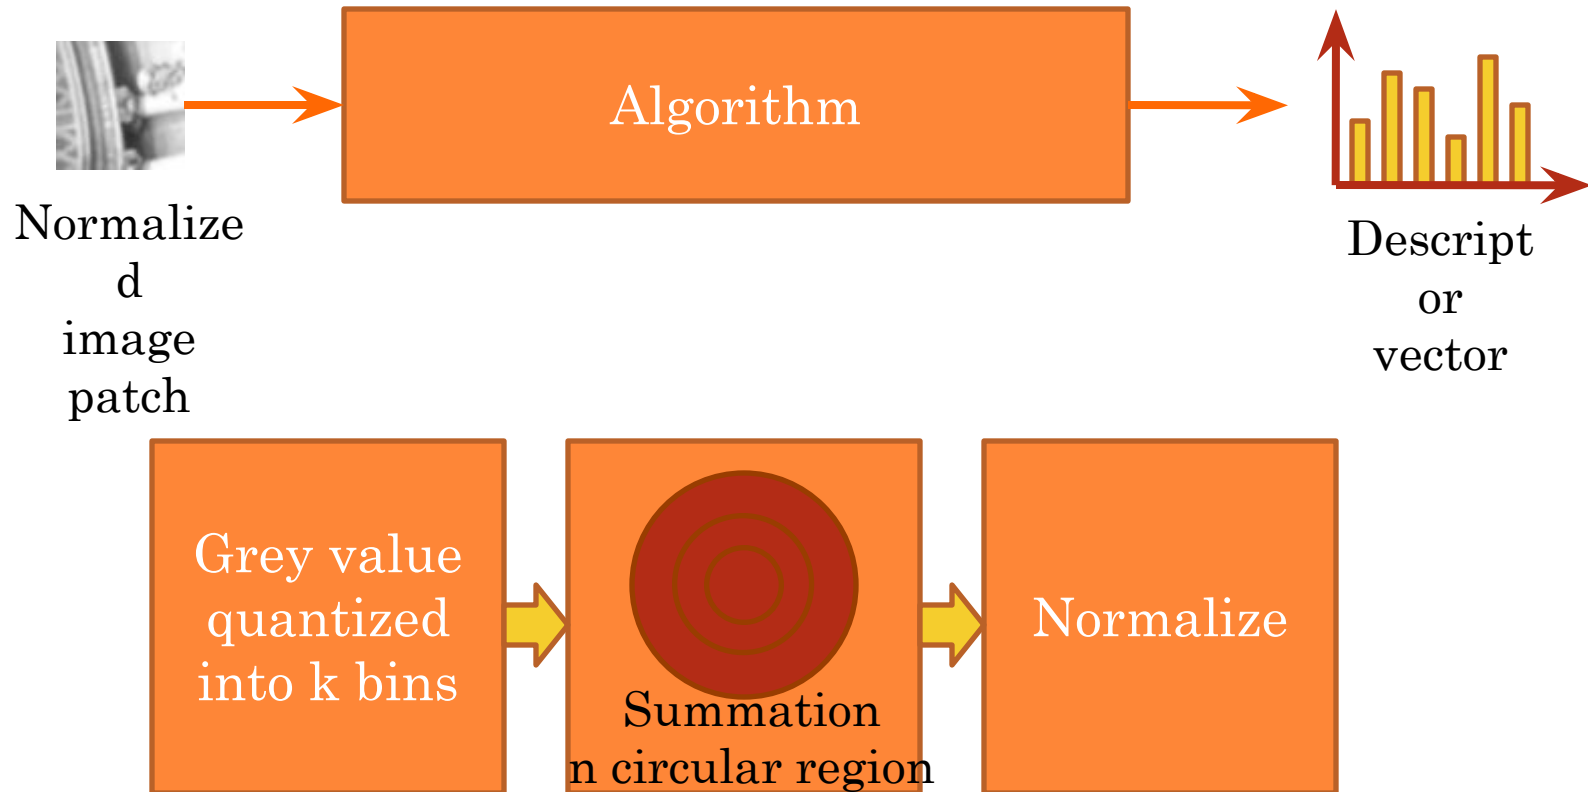
DESCRIPTOR ALGORITHMS



[Shape Context – Belongie et al, NIPS'00]



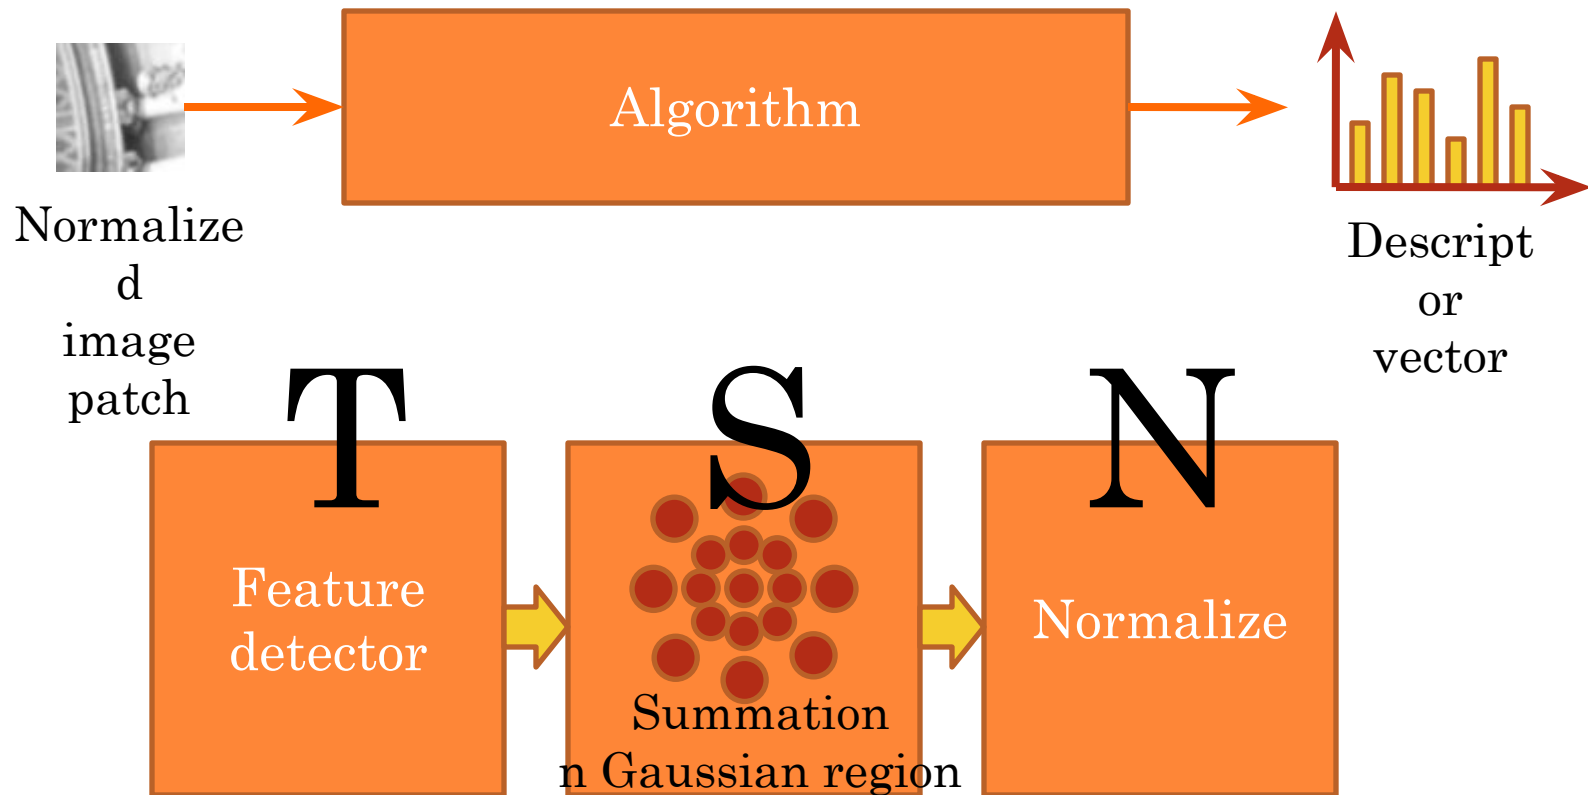
DESCRIPTOR ALGORITHMS



[Spin Image – Lazebnik et al, CVPR'03]



DESCRIPTOR ALGORITHMS




[Geometric Blur – Berg& Malik
CVPR'01]


[DAISY - Tola et al. CVPR'08]




OUR DESCRIPTOR ALGORITHM




S-BLOCK: PICKING THE BEST DAISY



□ S-Block: DAISY




T-BLOCK




□ T-Block:

- T1: Gradient bi-linearly weighted orientation binning
- T2: Rectified gradient $\{|\nabla_x| - \nabla_x, |\nabla_x| + \nabla_x, |\nabla_y| - \nabla_y, |\nabla_y| + \nabla_y\}$
- T3: Steerable filters




N-E-Q BLOCKS



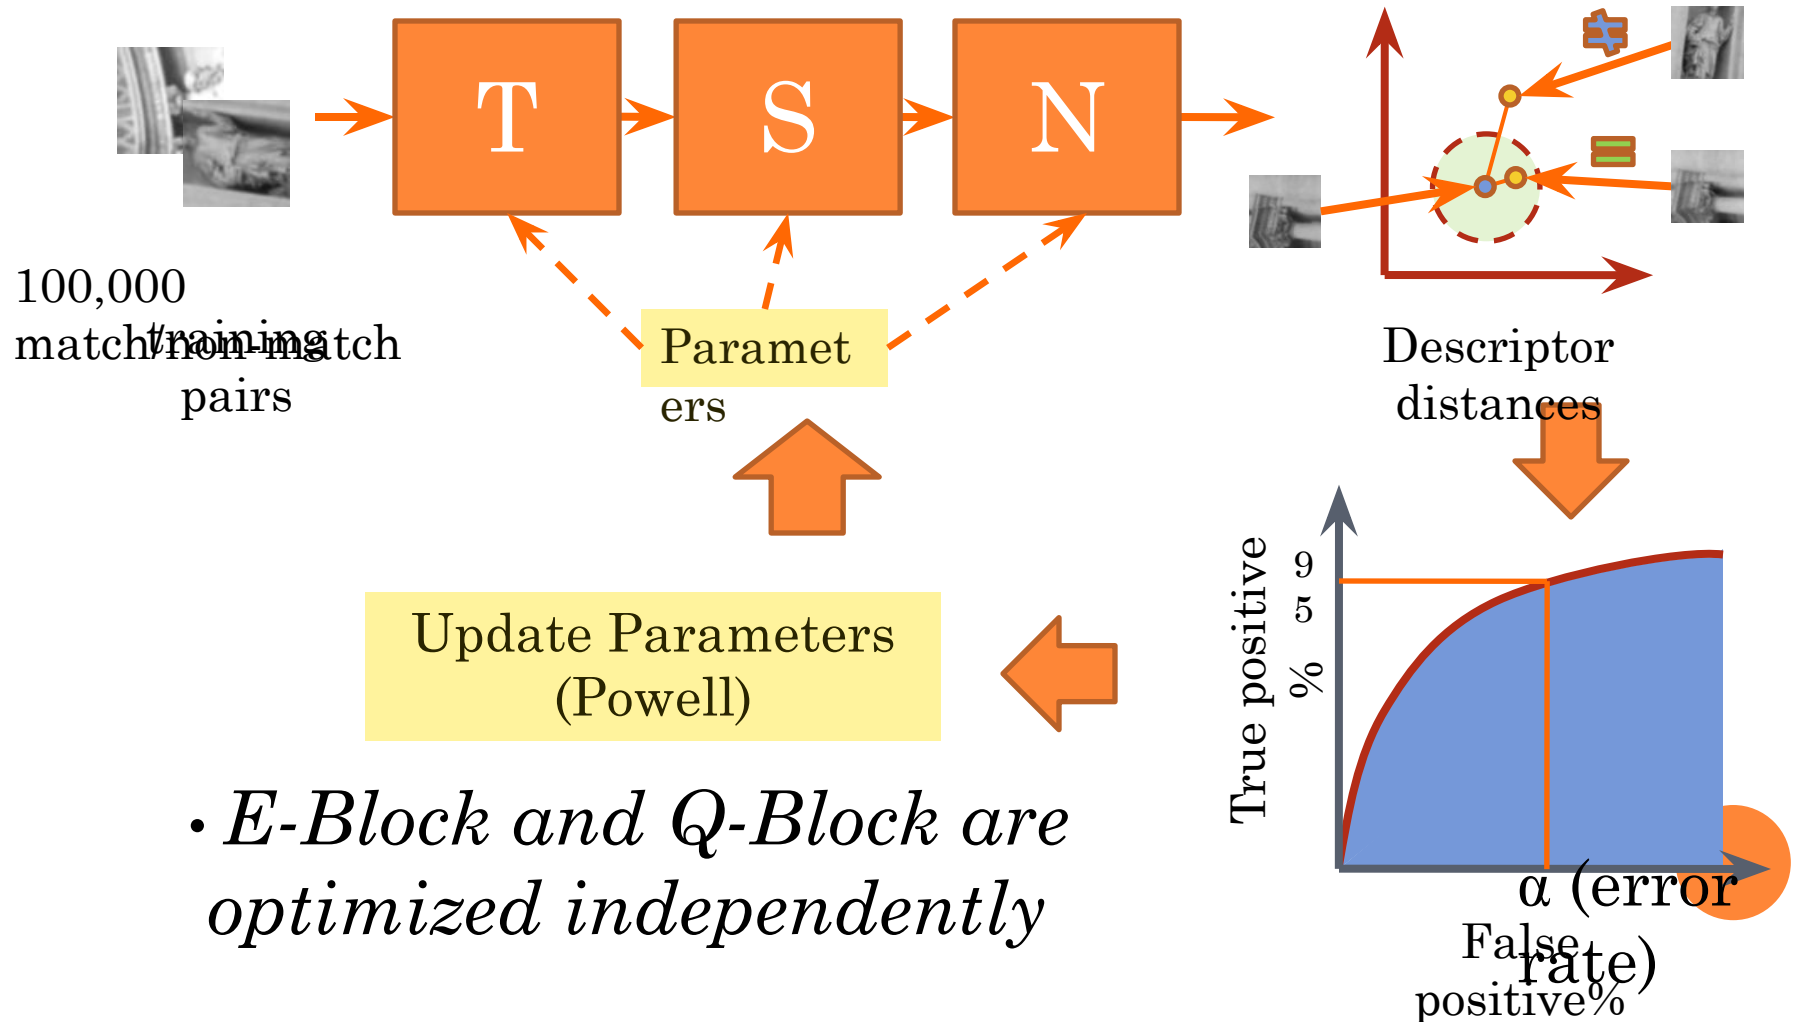
- N-Block: Uniform v.s. SIFT-like normalization Vector
- E-Block: Principle component analysis (PCA)
- Q-block: $q_i = \lfloor \beta L v_i \rfloor + \alpha$, β is learnt from data, $\alpha = 0.5$ if L is an odd number and $\alpha = 0$ otherwise.



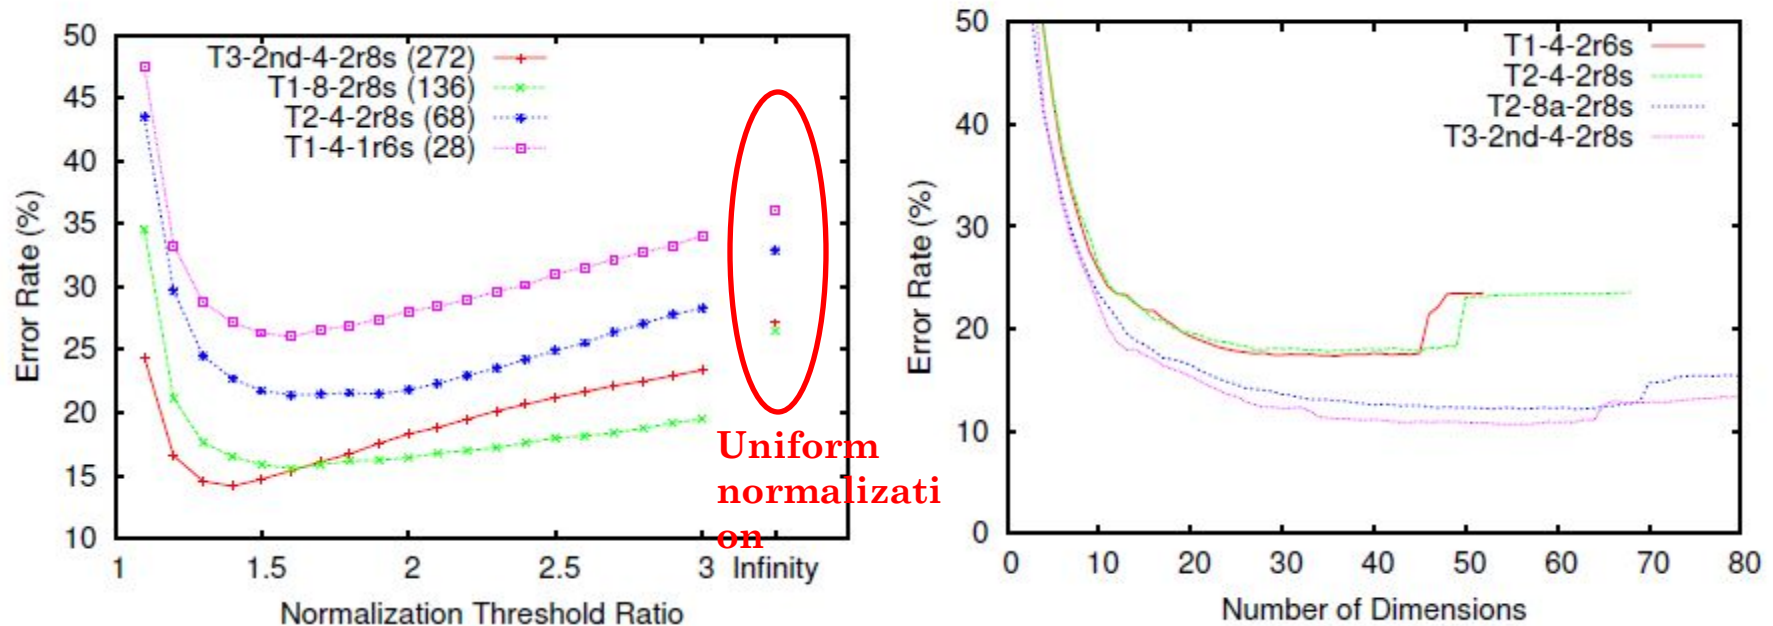
WHAT IS THE OPTIMAL CRITERION?



DISCRIMINATIVE LEARNING AND OPTIMAL CRITERION





SIFT-LIKE NORMALIZATION & PCA



- SIFT like normalization has a clear sweet spot.
- PCA can usually reduce the dimension to 30~50.


THE MOST DISCRIMINATIVE DAISY



- Steerable filters at two spatial scales with PCA
 - T3-2nd-6-2r6s + 2-scale + PCA: 42 dimension, 9.5%
 - $2 \times 6 \times 2 \times 2 \times (2 \times 6 + 1) = 624$ dimensions before PCA



THE MOST COMPACT DAISY

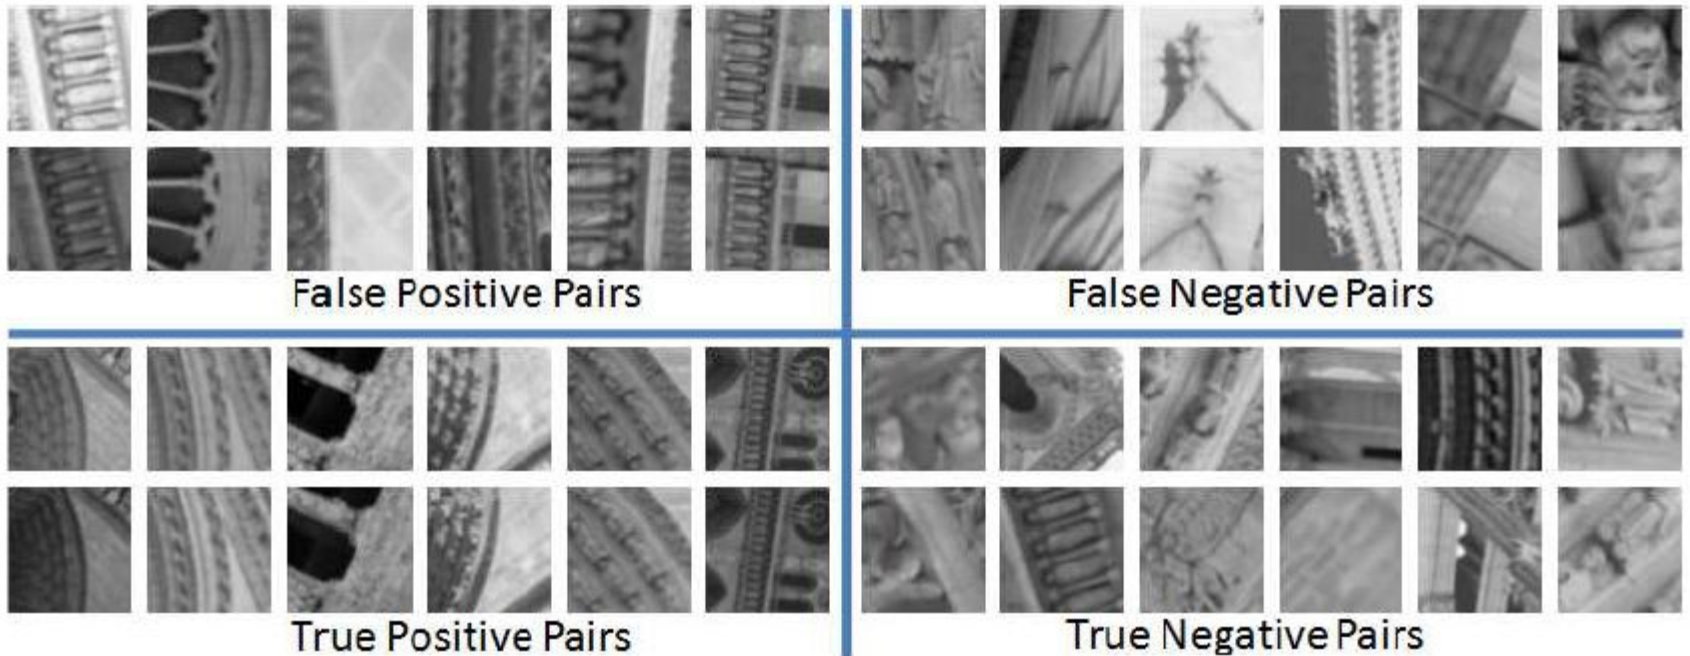


- **T3-2nd-4-2r8s**
 - 13 bytes, 13.2%
- **T2-4-1r8s-PCA**
 - 7.5 bytes, 24.4%
- **SIFT**
 - 128 bytes, 26.1%

- 2 bits /dimension is sufficient before PCA
- 4 bits /dimension is needed after PCA
- <http://www.cs.ubc.ca/~mbrown/patchdata/patchdata.html>



SOME OF THE ERRORS MADE



□ The world is repeating itself....



SUMMARY

□ Blob detection

- Brief of Gaussian filter
- Scale selection
- Lapacian of Gaussian (LoG) detector
- Difference of Gaussian (DoG) detector
- Affine co-variant region

□ Learning local descriptors

- How can we get ground-truth data?
- What is the form of the descriptor function?
- What is the optimal criterion?
- How do we optimize it?

